



Distributed file systems


Distributed file systems

- Introduction
- Clustering with Ceph
 - Types of Ceph client
- Preparation of Ceph cluster
 - Set Ceph nodes
 - Set a RDB client

Introduction

What are filesystems ?

- Organization of mass storage
 - To store data
 - To retrieve data
 - To share data between multiple programs
- Two points of view
 - User : a filename including pathname and its content
 - Operating system : metadata, pointers, blocks, chunks,...



Types of filesystems

- Local filesystems

- On one computer
- Direct access to a block device
- Managed by one OS
- XFS, BTRFS, EXT*, ...

- Shared filesystems

- Multiple OS access to one exported filesystem
- The filesystem is connected to one computer
- This computer manages its local block device
- NFS, CIFS



- But problems
 - I/O bottle necks
 - Network latency
 - Host is a SPOF

Solutions ?

- HA cluster


- One system active
- one system standby
- One cache

- Cluster filesystems

- One block device
- Block level access
- I/O cache consistency management : cache synchronization
- Multiple kernel handle I/O at any time
- SAN with OCFS, RedHat GFS, ...


A solution : Distributed Filesystems

- A file system spread over several nodes
- Access by multiple clients
 - Large and scalable capacities
 - High I/O throughput
 - Redundancy
- Transparency
 - Access
 - Location
 - Concurrency
 - Failure
- Examples
 - GlusterFS, Gfarm, Ceph,...



How it works

- **Multiple nodes**
 - They run their own OS
 - They access to one block device
- **One or more master nodes**
 - Dispatch requests
 - Manage filesystem status
 - Manage synchronization
- **Clients**
 - Mount locally logical filesystem
 - Served by different nodes



Ceph

What is CEPH ?

- Distributed storage system

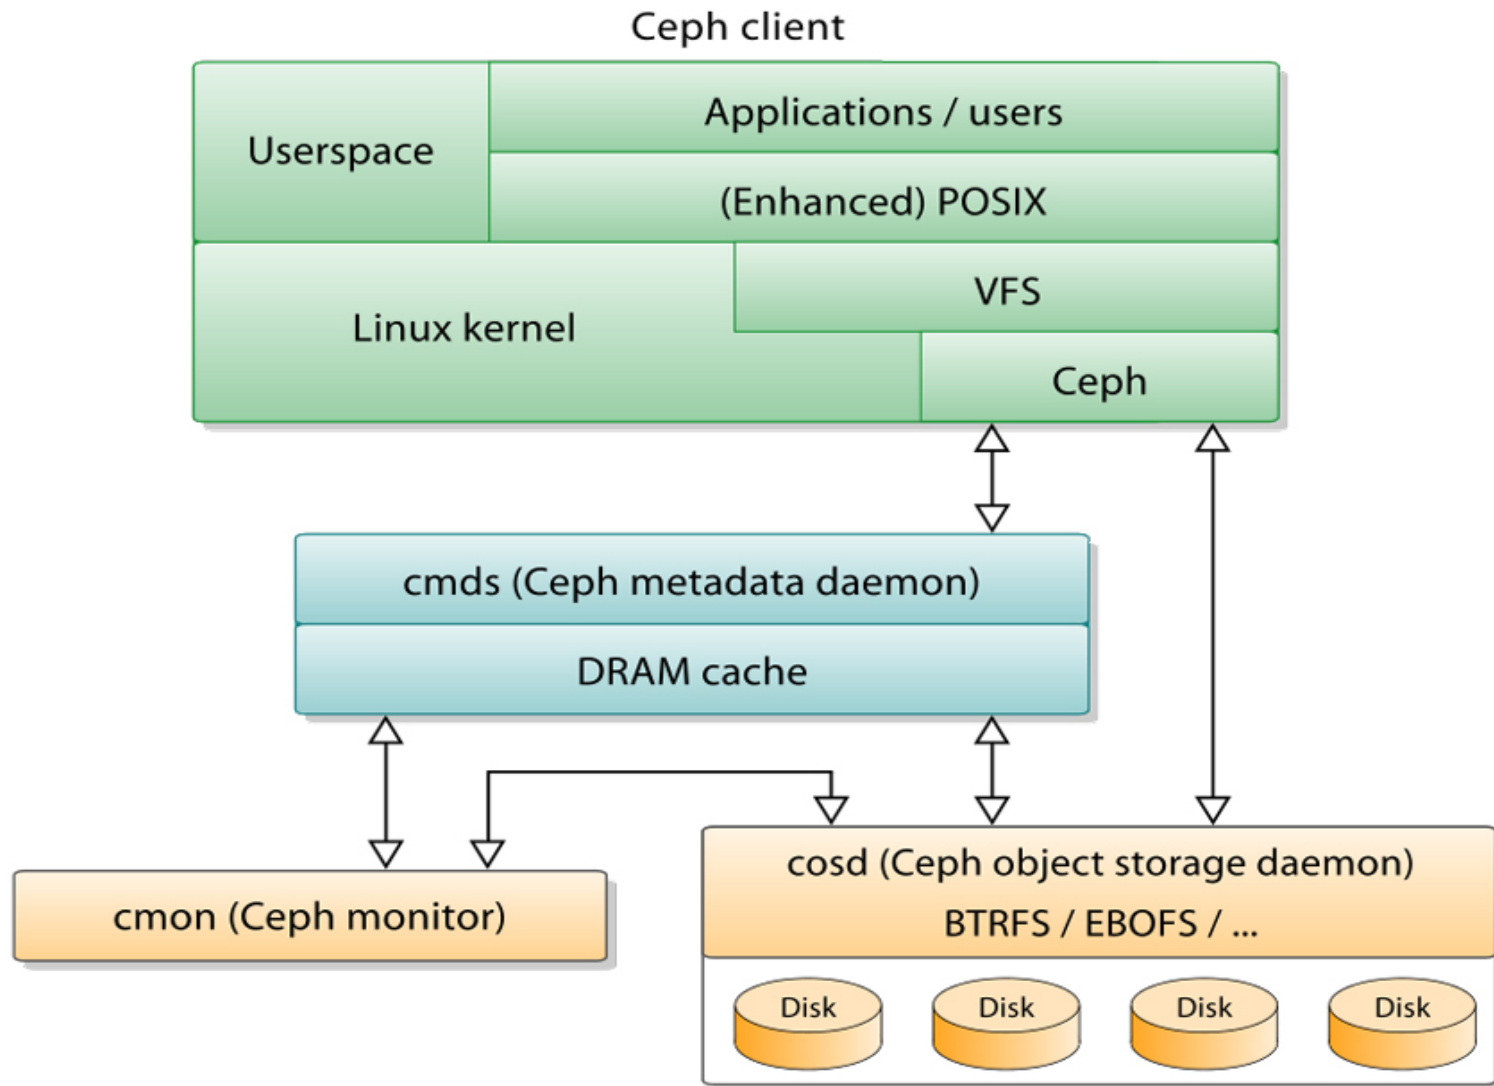
- Software-designed
- Stores **objects, blocks, files**
- Self-healing
- Scalable
- Designed for cloud infrastructure
- All components are redunded
 - No SPOF
 - multiactive

- Open Source

- Sage Weil, doctoral dissertation in 2006 (UCSC)
 - 2012 : Inktank Storage company
 - 2014 : Redhat aquires Inktank Storage
 - april 2016 : version 10.2.0


Organisation

- Several daemons to deliver service



Components : nodes

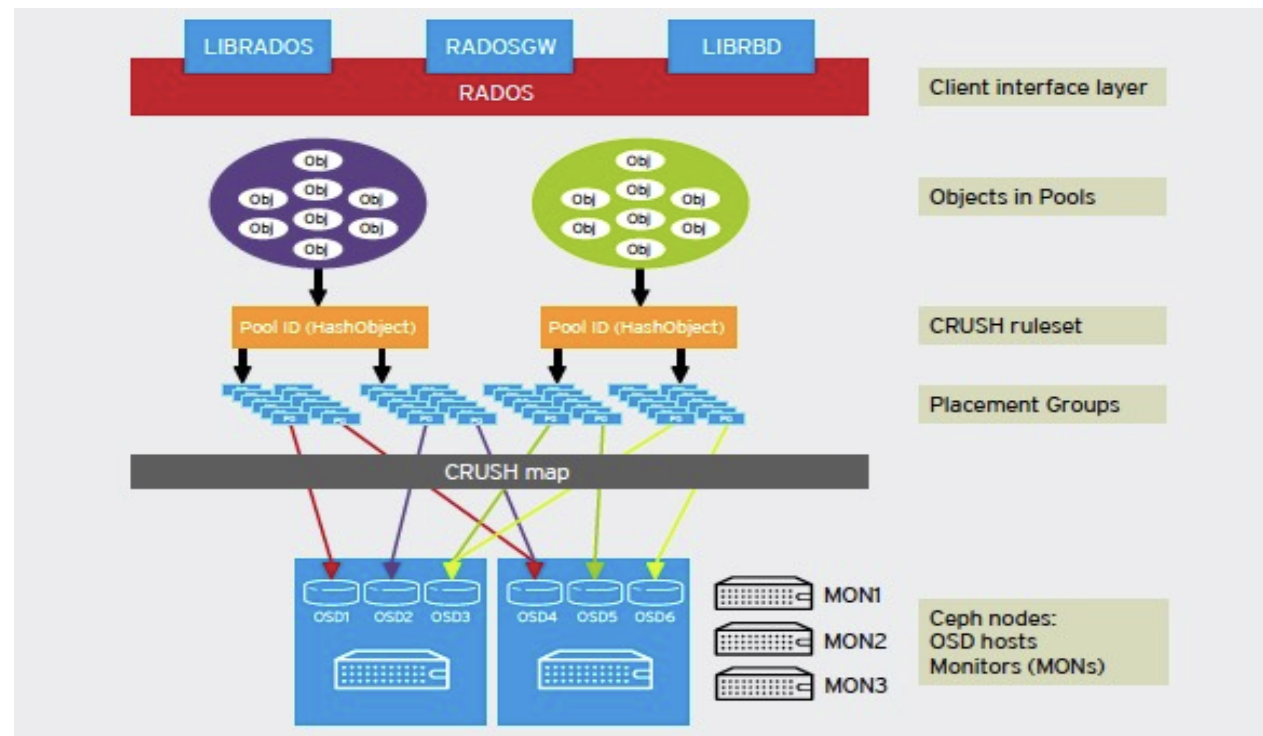
- An administration server : **admin**
 - Recommended for an easier deployment
 - For administrative task on the cluster
- A node running the metadata server : **mds**
 - Stores description of the user data
 - Load management
- At least one node running the monitor : **mon**
 - Clients get a copy of the 5 cluster maps
 - Monitor map (id, address, port,...)
 - OSD map (id, data containers, status, ...)
 - PG – Placement group - map (status, ...)
 - CRUSH map (list of storage devices, rules, ...)
 - MDS map (list of mds, ...)
 - Clusterized monitors
- At least one node running the Object Storage Device : **osd**
 - Store objects on disks
 - Reports to the **mon**




Architecture : another view

- **Distribution and Modularity**

- Pools : ownership/access to objects, replication rules, number of PG
- PG : logical container used to spread objects on OSDs (balancing)
- CRUSH ruleset : clients compute data location and access directly to data



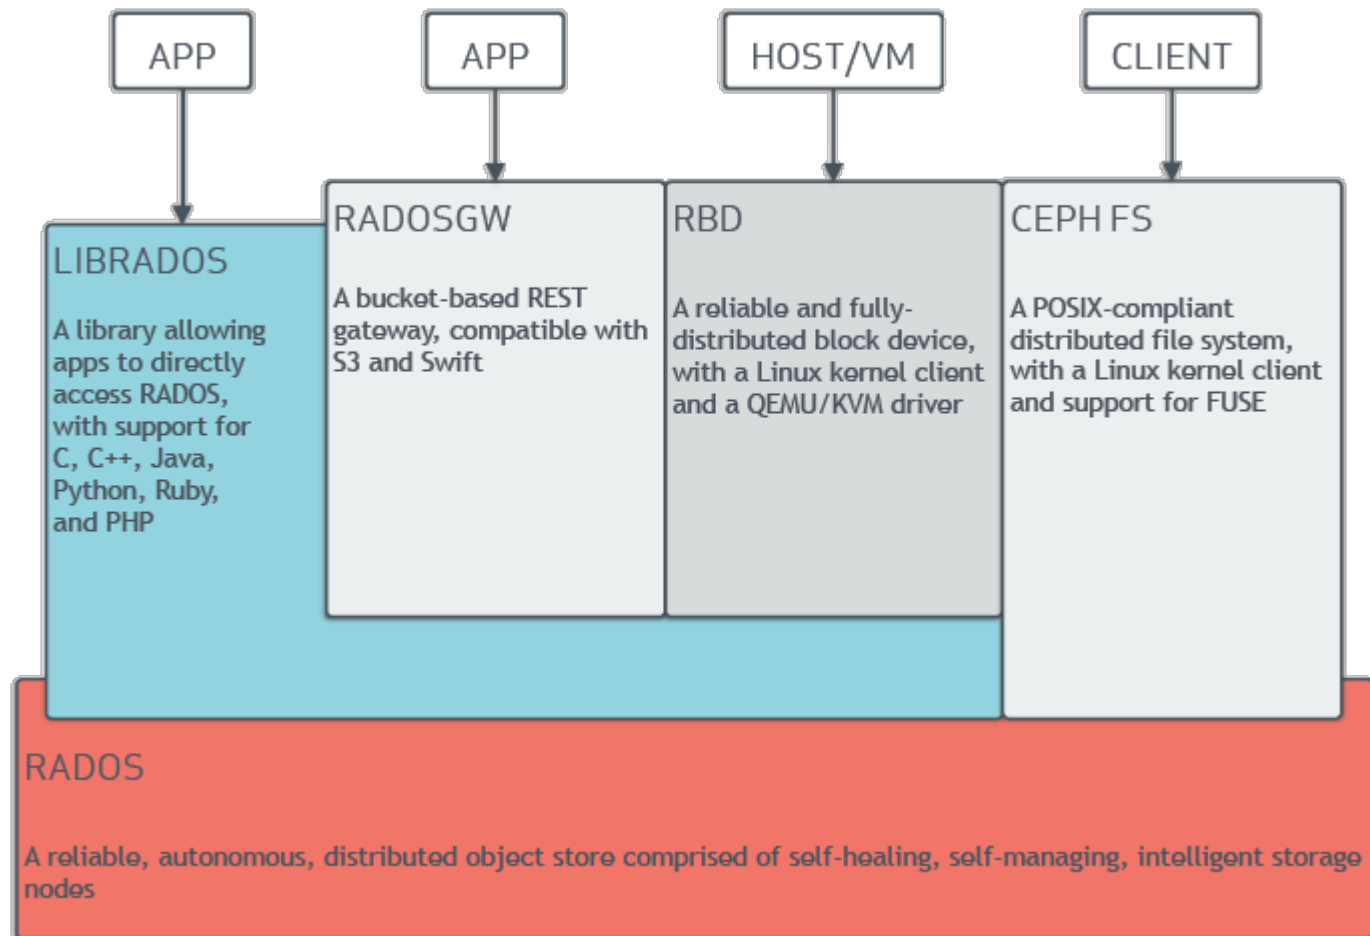
- Example with a write access
 - 1. client gets the updated cluster map from a MON server
 - 2. data are converted in objects (with object ID and pool ID)
 - 3. CRUSH determines PG and primary OSD where to store objects
 - 4. client contacts primary OSD node to store objects
 - 5. with CRUSH, primary OSD node searches secondary Pgs and secondary OSD nodes
 - 6. primary OSD replicates objects on secondary OSD nodes



Types of client

Clients

- With libraries, OS or client software, client nodes access to
 - Objects
 - Blocks
 - Files



Data storage

- OSD nodes

- store objects
- On a flat namespace
- No hierarchy

- An object has

- An ID
- Binary data
- Metadata (name/value pairs)
- The client puts the semantic it wants
 - Blocks
 - Objects
 - Files

ID	Binary Data	Metadata	
1234	0101010101010100110101010010 0101100001010100110101010010 0101100001010100110101010010	name1 name2 nameN	value1 value2 valueN

- Clients and OSD nodes compute data location
 - Hash algorithm
 - CRUSH : Controlled Replication Under Scalable Hashing
 - CRUSH used by clients to store objects
 - CRUSH used by OSD nodes to store objects and their replicas
 - Maps are updated

```
root default
  datacenter loi
    room loire-presidenc
      rack karuizawa
        host ceph-osd-loi-A-1-1
          osd.12 up
      rack hazelburn
        host ceph-osd-loi-A-2-1
          osd.15 up
  datacenter lmb
    room lombarderie-ltc
      rack kavalan
        host ceph-osd-lmb-A-1-1
          osd.6 up
```

- RADOS
 - Lowest layer of Ceph storage clusters
 - Reliable Autonomic Distributed Object Store



- Used to develop Ceph Clients
 - For a parallel access object storage
 - Asynchronous communication protocol
 - Actions
 - R/W
 - Create/remove
 - Append/truncate
 - Manage XATTR's
 - Manage Name/Value pairs
 - Data stripping
 - ...
 - Library **librados**



RBD : RADOS Block Device

- Block device service

- Clients access

- with Xen, QEMU or libvirt (**qemu-img create -f rbd rbd:image 2G**)
- Kernel Ceph block device (**modprobe rbd**)
- Command **rbd**

- Stores blocks over multiple objects on Ceph

- Objects are mapped on PG (placement group)

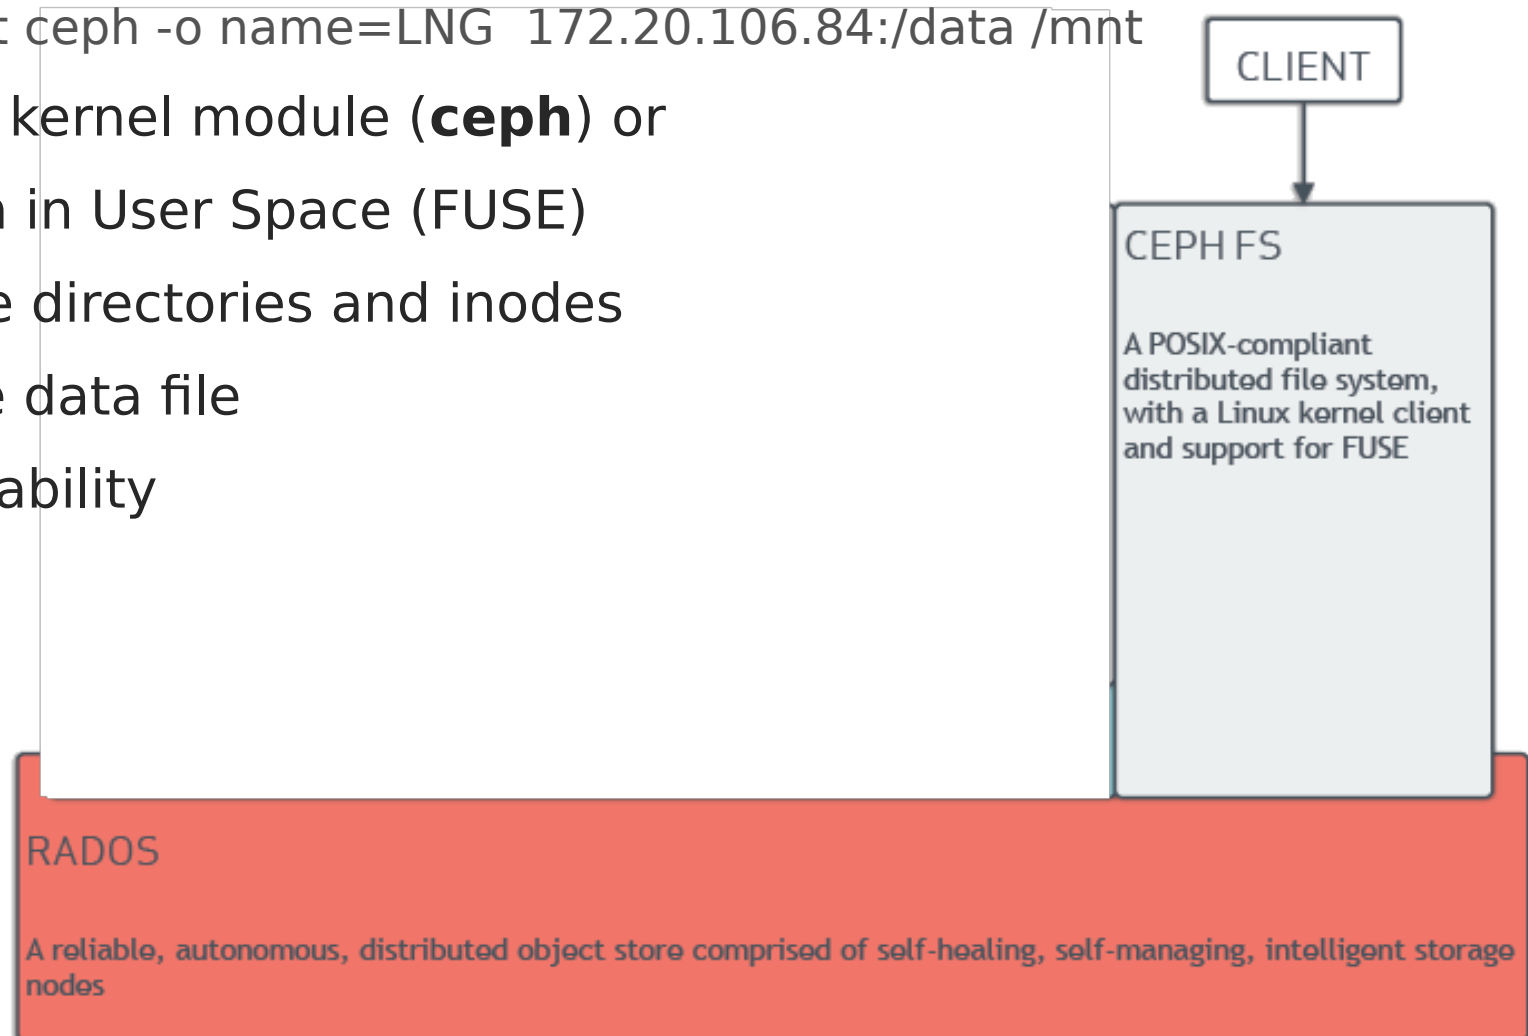
- Choice of FS format (BTRFS is recommended)



CephFS : Distributed File System

- CephFS

- Clients access with **mount** command
 - `mount -t ceph -o name=LNG 172.20.106.84:/data /mnt`
- Client use kernel module (**ceph**) or
- Filesystem in User Space (FUSE)
- **MDS** store directories and inodes
- **OSD** store data file
- High-availability
- Scalability



RADOS Gateway

- RESTful interface
 - Access to objects only
 - Two standard APIs used by apps based on clusters
 - Amazon-S3
 - OpenStack-Swift
 - Based on Ceph Object Gateway daemon : **radosgw**



Preparation

Needs

- For a test cluster : 5 nodes
 - cluster-admin
 - cluster-mon
 - cluster-mds
 - cluster-osd1
 - cluster-osd2
- For a production cluster : 10 nodes
 - cluster-admin
 - cluster-mon1, cluster-mon1, cluster-mon2
 - cluster-mds1, cluster-mds2
 - cluster-osd1, cluster-osd2, cluster-osd3, cluster-osd4,
- Minimum hardware
 - 1 CPU , 1 GB RAM per node
- User
 - A linux user « **ceph** » on all nodes

- SSH: distribute the **ceph** user SSH-key on all nodes
- Sudo : give root access to ceph user
- Name resolution
 - Give a significant hostname to nodes
 - Set DNS records for all
- Prepare control node : on ceph-admin
 - Install ceph-deploy package
- Set role to monitor node(s)
 - Ceph-deploy new cluster-mon
- Install Ceph on nodes
 - ceph-deploy install cluster-admin cluster-mon cluster-mds cluster-osd1 cluster-osd2

- Define the monitor node
 - `ceph-deploy mon create cluster-mon`
- Get authentication keys from the monitor
 - `ceph-deploy gatherkeys cluster-mon`
- Select disks on data nodes (ie : `/dev/sdb`)
 - `ceph-deploy disk zap cluster-osd1:sdb` (same for `osd2`)
- Prepare data nodes
 - `ceph-deploy osd prepare cluster-osd1:sdb` (same for `osd2`)
- Enable data nodes
 - `ceph-deploy osd activate cluster-osd1:sdb` (same for `osd2`)
- Define MDS nodes
 - `ceph-deploy mds create cluster-mds`

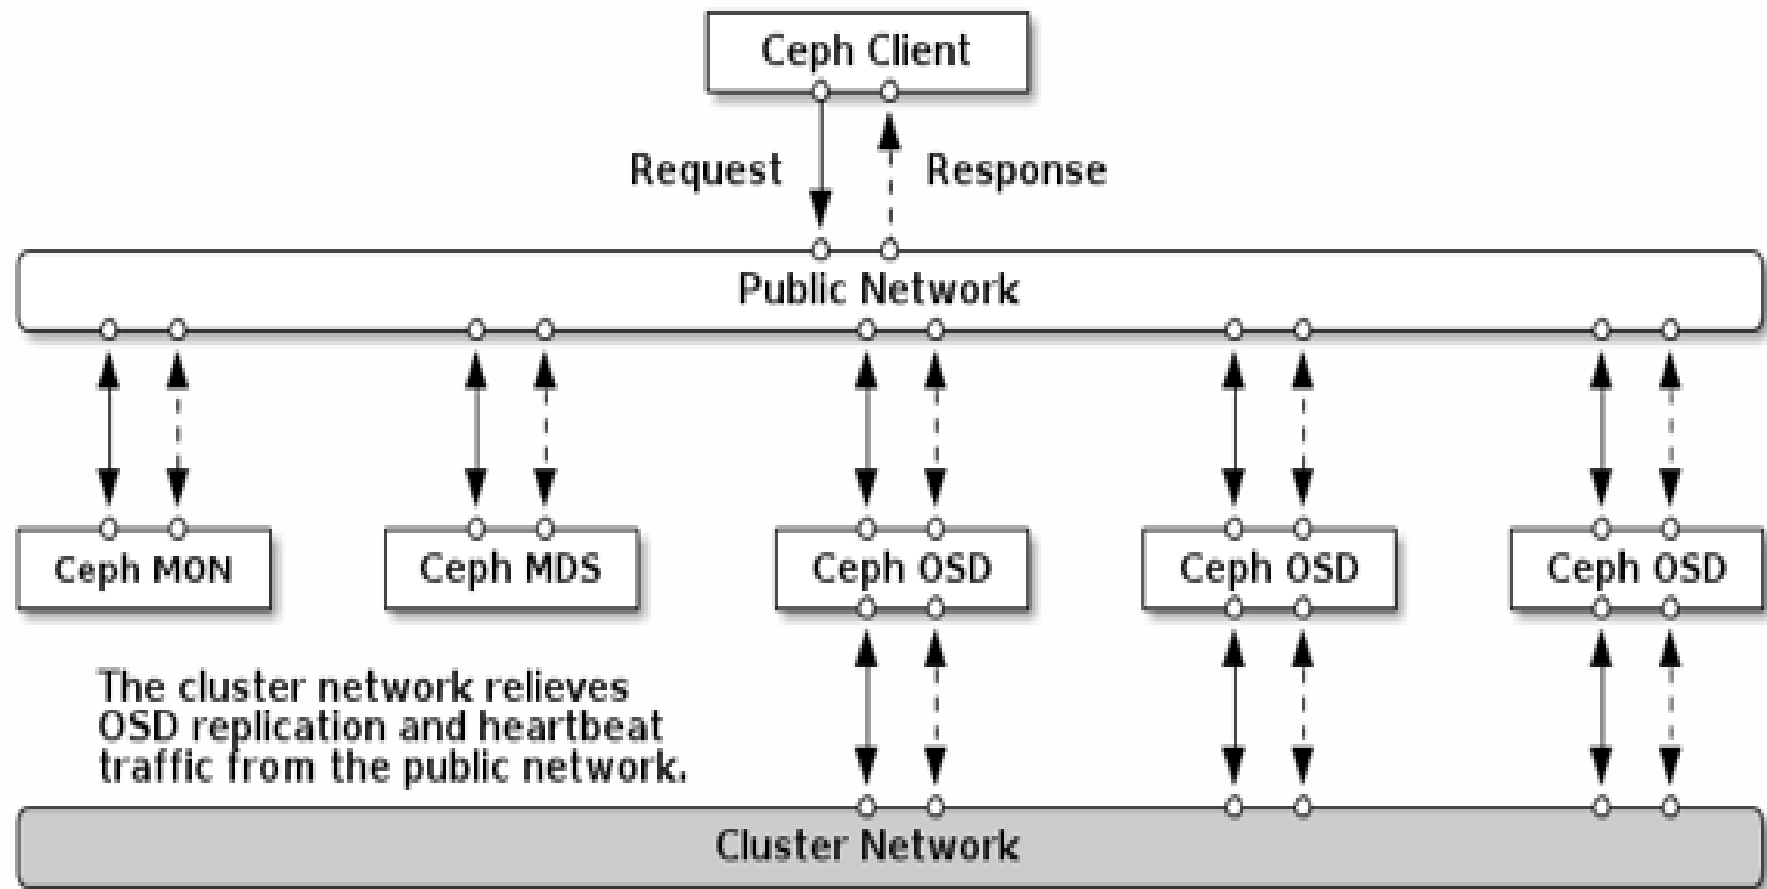
It works ?

- Check it works
 - ceph status

```
$ ceph status
cluster a5bf6b38-02f4-43fb-bffc-9c932e98888e
health HEALTH_OK
```

- Check it works
 - ceph health

Example with 3 OSD nodes



RDB client access

Create a client

- On client node called « cluster-client1 »
 - Create Linux « ceph » user
 - Install SSH keys
 - Set Sudo access
- Set DNS records for the new node
- On cluster-admin node
 - `ceph-deploy install cluster-client1`

Create block devices

- Create a pool called « volumes » containing objects
 - `rados mkpool volumes`
- Create two objects (1GB each)
 - `rdb -p volumes - -size 1024 create vol1`
 - `rdb -p volumes - -size 1024 create vol2`
- Map new objects (vol1, vol2) with block devices
 - `rbd map volumes/vol1`
 - `rbd map volumes/vol2`
- Check new blocks devices exist
 - `/dev/rdb/volumes/rdb0`
 - `/dev/rdb/volumes/rdb1`
- Format and use them !

Thanks for your attention

LINAGORA - headquarters

80, rue Roque de Fillol
92800 PUTEAUX
FRANCE

Phone : +33 (0)1 46 96 63 63

Info : info@linagora.com

Web : www.linagora.com

 **@linagora**

 **facebook.com/Linagora/**

